

```

# Data handling
import pandas as pd          # For loading and manipulating dataset
import numpy as np            # For numerical operations

# Text preprocessing
import re                      # For removing punctuation using regex
import string                  # For handling string operations

# Stopwords
from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS

# Feature extraction
from sklearn.feature_extraction.text import TfidfVectorizer # Convert text to numerical features

# Train-test split
from sklearn.model_selection import train_test_split

# Naive Bayes model
from sklearn.naive_bayes import MultinomialNB

# Evaluation metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

```

```

# Load dataset (Change path if needed)
# Assuming the CSV is comma-separated and has a header.
# The actual text content appears to be in a column named 'text'.
# The actual label content appears to be in a column named 'label'.
data = pd.read_csv("/content/news.csv", sep=',')

# Select the 'text' and 'label' columns and rename 'text' to 'message'
data = data[['text', 'label']].rename(columns={'text': 'message'})

# Display first 5 rows
print("First 5 Samples:\n")
print(data.head())

# Dataset size
print("\nDataset Shape:", data.shape)

# Class distribution
print("\nClass Distribution:\n")
print(data['label'].value_counts())

```

First 5 Samples:

	message	label
0	Daniel Greenfield, a Shillman Journalism Fellow...	FAKE
1	Google Pinterest Digg LinkedIn Reddit StumbleUpon...	FAKE
2	U.S. Secretary of State John F. Kerry said Monday...	REAL
3	— Kaydee King (@KaydeeKing) November 9, 2016 They...	FAKE
4	It's primary day in New York and front-runners...	REAL

Dataset Shape: (6335, 2)

## Class Distribution:

```
label
REAL    3171
FAKE    3164
Name: count, dtype: int64
```

```
# Function to clean text
def preprocess_text(text):

    # Handle non-string inputs (e.g., NaN values)
    if not isinstance(text, str):
        return ""

    # Convert to lowercase
    text = text.lower()

    # Remove punctuation
    text = re.sub(f"[{string.punctuation}]", "", text)

    # Remove stopwords
    words = text.split()
    words = [word for word in words if word not in ENGLISH_STOP_WORDS]

    return " ".join(words)

# Apply preprocessing
data['clean_message'] = data['message'].apply(preprocess_text)

print("\nCleaned Sample:\n")
print(data[['message','clean_message']].head())

print("\nDataFrame Info after adding 'clean_message' column:\n")
data.info()
```

## Cleaned Sample:

```
          message \
0 Daniel Greenfield, a Shillman Journalism Fello...
1 Google Pinterest Digg Linkedin Reddit Stumbleu...
2 U.S. Secretary of State John F. Kerry said Mon...
3 – Kaydee King (@KaydeeKing) November 9, 2016 T...
4 It's primary day in New York and front-runners...

          clean_message
0 daniel greenfield shillman journalism fellow f...
1 google pinterest digg linkedin reddit stumbleu...
2 secretary state john f kerry said monday stop ...
3 – kaydee king kaydeeking november 9 2016 lesso...
4 primary day new york frontrunners hillary clin...
```

## DataFrame Info after adding 'clean\_message' column:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6335 entries, 0 to 6334
Data columns (total 3 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   message     6335 non-null   object 
 1   label       6335 non-null   object 
 2   clean_message 6335 non-null  object 
dtypes: object(3)
```

memory usage: 148.6+ KB

```
# Initialize TF-IDF vectorizer
vectorizer = TfidfVectorizer()

# Convert text into numerical feature matrix
X = vectorizer.fit_transform(data['clean_message'])

# Target variable
y = data['label']

print("Feature Matrix Shape:", X.shape)
print("Sample Feature Names:", vectorizer.get_feature_names_out()[:10])

Feature Matrix Shape: (6335, 84068)
Sample Feature Names: ['00' '000' '0000' '0000000031' '000000031' '0000035' '00001400' '00006'
 '00011' '00017b2908ff9fa45188d243fd49aaeeb2dhrcofficecom']
```

```
# Split data into training and testing sets
# Using 80% for training and 20% for testing, with a random state for reproducibility
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)

X_train shape: (5068, 84068)
X_test shape: (1267, 84068)
y_train shape: (5068,)
y_test shape: (1267,)
```

```
# Initialize Multinomial Naive Bayes
model = MultinomialNB()

# Train model
model.fit(X_train, y_train)

print("Model Parameters:\n", model.get_params())

Model Parameters:
{'alpha': 1.0, 'class_prior': None, 'fit_prior': True, 'force_alpha': True}
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV

# Define the parameter grid for Logistic Regression
# Using 'liblinear' solver for 'l1' penalty which is good for small datasets
param_grid_log_reg = {
    'C': [0.1, 1, 10, 100],
    'penalty': ['l1', 'l2']
}

# Initialize Logistic Regression model
# Set solver to 'liblinear' as it supports both 'l1' and 'l2' penalties
# and is good for relatively small datasets.
log_reg = LogisticRegression(solver='liblinear', random_state=42, max_iter=1000)

# Initialize GridSearchCV
grid_search_log_reg = GridSearchCV(log_reg, param_grid_log_reg, cv=5, scoring='accuracy', verbose
```

```
# Fit GridSearchCV to the training data
grid_search_log_reg.fit(X_train, y_train)

# Print the best parameters and best score
print("Best parameters for Logistic Regression:", grid_search_log_reg.best_params_)
print("Best cross-validation accuracy for Logistic Regression:", grid_search_log_reg.best_score_)

Fitting 5 folds for each of 8 candidates, totalling 40 fits
Best parameters for Logistic Regression: {'C': 100, 'penalty': 'l2'}
Best cross-validation accuracy for Logistic Regression: 0.9327170842168184
```

```
best_log_reg_model = grid_search_log_reg.best_estimator_
y_pred_tuned_log_reg = best_log_reg_model.predict(X_test)

accuracy_tuned_log_reg = accuracy_score(y_test, y_pred_tuned_log_reg)
precision_tuned_log_reg = precision_score(y_test, y_pred_tuned_log_reg, pos_label="FAKE")
recall_tuned_log_reg = recall_score(y_test, y_pred_tuned_log_reg, pos_label="FAKE")
f1_tuned_log_reg = f1_score(y_test, y_pred_tuned_log_reg, pos_label="FAKE")

print("Accuracy (Tuned Logistic Regression):", accuracy_tuned_log_reg)
print("Precision (Tuned Logistic Regression):", precision_tuned_log_reg)
print("Recall (Tuned Logistic Regression):", recall_tuned_log_reg)
print("F1-Score (Tuned Logistic Regression):", f1_tuned_log_reg)

print("\nClassification Report (Tuned Logistic Regression):\n")
print(classification_report(y_test, y_pred_tuned_log_reg))

cm_tuned_log_reg = confusion_matrix(y_test, y_pred_tuned_log_reg)

plt.figure(figsize=(5,4))
sns.heatmap(cm_tuned_log_reg, annot=True, fmt='d', cmap='Blues',
            xticklabels=best_log_reg_model.classes_,
            yticklabels=best_log_reg_model.classes_)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix for Tuned Logistic Regression")
plt.show()
```

Accuracy (Tuned Logistic Regression): 0.9400157853196527  
Precision (Tuned Logistic Regression): 0.9259259259259259  
Recall (Tuned Logistic Regression): 0.9554140127388535  
F1-Score (Tuned Logistic Regression): 0.9404388714733543

Classification Report (Tuned Logistic Regression):

	precision	recall	f1-score	support
FAKE	0.93	0.96	0.94	628
REAL	0.95	0.92	0.94	639
accuracy			0.94	1267
macro avg	0.94	0.94	0.94	1267
weighted avg	0.94	0.94	0.94	1267

Confusion Matrix for Tuned Logistic Regression

